SUPRIYA S

Florida USA • (904)-274-2321 • sup96478@gmail.com

SUMMARY

Data Engineer with 5+ years building distributed data-processing systems and analytics-ready data models on AWS and Databricks. Strong in Python, PySpark, Scala, and SQL, with Spark/Kafka streaming experience and hands-on Flink exposure, applied to web-scale datasets. Proven record of cutting pipeline latency and Spark runtimes, owning data quality, lineage, and documentation, and translating ambiguous requirements into reliable pipelines that power analytics and ML. Collaborative across engineering, analytics, and data science; takes intelligent risks in fast-changing environments.

TECHNICAL SKILLS

Languages: Python, PySpark, Scala, SQL, Java (familiar)
Frameworks/Engines: Apache Spark, Spark Streaming, Kafka, Flink (familiar)
Cloud & Platforms: AWS (S3, EMR, Glue, Redshift, RDS, CloudWatch), Databricks, Snowflake
Data Stores: HDFS, Cassandra, NoSQL, Relational DBs
DevOps & CI/CD: Jenkins, Git, Groovy, shell scripting
Data Modeling & Quality: Dimensional modeling, normalization/denormalization, data testing, monitoring, lineage
Other: ETL, streaming, batch processing, schema design, performance tuning, observability, ML feature pipelines

EXPERIENCE

Data Engineer Sep 2022 – Nov 2024
Accenture *Florida, USA*

- Cut data-processing latency 40% by architecting distributed batch and real-time ETL pipelines on AWS Glue, PySpark, and Databricks that ingested web-scale data from 12+ heterogeneous sources into Amazon S3 data lakes.
- Improved pipeline stability ~25% by re-implementing performance-critical Spark jobs in Scala for stronger type safety and JVM performance on web-scale datasets.
- Enabled real-time analytics at ~50K events/sec by integrating Kafka with Spark Streaming for low-latency ingestion into Cassandra, powering [product feature / use case].
- Reduced a critical Spark job's runtime from ~2 hours to ~45 minutes by profiling execution plans and tuning partitioning, cluster config, and UDFs after code-review feedback.
- Lowered warehouse spend and query time ~25% by designing dimensional models (normalization/denormalization) and authoring Spark SQL transformations across Snowflake and Redshift.
- Improved data reliability by building data-quality checks, monitoring, and lineage into pipelines, cutting data defects ~30%.
- Accelerated releases from biweekly to near-daily by automating CI/CD with Jenkins and Groovy.
- Delivered 20+ documented, reusable pipelines by translating ambiguous requirements with data scientists, analysts, and business stakeholders.

Application Development Analyst May 2017 – Apr 2020
Accenture *Hyderabad, India*

- Standardized ingestion across 6+ projects by building reusable PySpark modules and automated transforms, improving developer velocity ~30%.
- Authored Scala-based Spark transformations for complex business logic, improving processing efficiency ~20% over equivalent Python/SQL implementations.
- Reduced reporting query time ~35% by implementing and optimizing SQL pipelines in Snowflake and Redshift.

- Cut incident MTTR ~40% by building monitoring dashboards and structured root-cause analyses.
- Built backend ETL services moving data from mainframe and transactional systems into HDFS and Cassandra to enable analytics outcome.

CERTIFICATIONS

- AWS Certified DevOps Engineer – Professional, Amazon Web Services
- AWS Certified Generative AI Developer – Professional, Amazon Web Services

EDUCATION

Bachelors in Electronics and Communication Engineering 2013 - 2017